

Technical assessment

October 23, 2018

1 Bicycle sharing demand Prediction

This is a city bicycle rented system I'm provided Washington DC bicycle rented records per hour in two years, train datasets include every month first 19 days and test datasets consist of last 10 days (we need to predict this part of time period).

2 Data load and Analysis

we are going to use pandas in python to do data analysis ** numpy is also indispensable **

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import os

In [56]: Folder_Path = '/Users/songzhewei/Desktop/technical assessment'
SaveFile_Name = 'DC2011-2012.csv'
def read(path, newFileName):
    files = os.listdir(path)
    with open(path + "/" + newFileName, "w") as f:
        for file in files:
            if file != newFileName:
                with open(path + "/" + file) as f1:
                    while 1:
                        line = f1.readline()
                        if not line:
                            break
                        f.write(line)
                    f.write("\n")

In [ ]: df=read(Folder_Path, SaveFile_Name)

In [58]: df1 = pd.read_csv("2011.csv")
df2 = pd.read_csv("2012.csv")

In [59]: df=df1.append(df2)
```

load data into cacheshow it firstlet's see first 10 rows

In [60]: df

```
Out[60]:
```

	datetime	season	holiday	workingday	weather	temp	atemp	\
0	2011/1/1 0:00	1	0	0	1	9.84	14.395	
1	2011/1/1 1:00	1	0	0	1	9.02	13.635	
2	2011/1/1 2:00	1	0	0	1	9.02	13.635	
3	2011/1/1 3:00	1	0	0	1	9.84	14.395	
4	2011/1/1 4:00	1	0	0	1	9.84	14.395	
5	2011/1/1 5:00	1	0	0	2	9.84	12.880	
6	2011/1/1 6:00	1	0	0	1	9.02	13.635	
7	2011/1/1 7:00	1	0	0	1	8.20	12.880	
8	2011/1/1 8:00	1	0	0	1	9.84	14.395	
9	2011/1/1 9:00	1	0	0	1	13.12	17.425	
10	2011/1/1 10:00	1	0	0	1	15.58	19.695	
11	2011/1/1 11:00	1	0	0	1	14.76	16.665	
12	2011/1/1 12:00	1	0	0	1	17.22	21.210	
13	2011/1/1 13:00	1	0	0	2	18.86	22.725	
14	2011/1/1 14:00	1	0	0	2	18.86	22.725	
15	2011/1/1 15:00	1	0	0	2	18.04	21.970	
16	2011/1/1 16:00	1	0	0	2	17.22	21.210	
17	2011/1/1 17:00	1	0	0	2	18.04	21.970	
18	2011/1/1 18:00	1	0	0	3	17.22	21.210	
19	2011/1/1 19:00	1	0	0	3	17.22	21.210	
20	2011/1/1 20:00	1	0	0	2	16.40	20.455	
21	2011/1/1 21:00	1	0	0	2	16.40	20.455	
22	2011/1/1 22:00	1	0	0	2	16.40	20.455	
23	2011/1/1 23:00	1	0	0	2	18.86	22.725	
24	2011/1/2 0:00	1	0	0	2	18.86	22.725	
25	2011/1/2 1:00	1	0	0	2	18.04	21.970	
26	2011/1/2 2:00	1	0	0	2	17.22	21.210	
27	2011/1/2 3:00	1	0	0	2	18.86	22.725	
28	2011/1/2 4:00	1	0	0	2	18.86	22.725	
29	2011/1/2 6:00	1	0	0	3	17.22	21.210	
...
5434	2012/12/18 18:00	4	0	1	1	15.58	19.695	
5435	2012/12/18 19:00	4	0	1	1	15.58	19.695	
5436	2012/12/18 20:00	4	0	1	1	14.76	16.665	
5437	2012/12/18 21:00	4	0	1	1	14.76	17.425	
5438	2012/12/18 22:00	4	0	1	1	13.94	16.665	
5439	2012/12/18 23:00	4	0	1	1	13.94	17.425	
5440	2012/12/19 0:00	4	0	1	1	12.30	15.910	
5441	2012/12/19 1:00	4	0	1	1	12.30	15.910	
5442	2012/12/19 2:00	4	0	1	1	11.48	15.150	
5443	2012/12/19 3:00	4	0	1	1	10.66	13.635	
5444	2012/12/19 4:00	4	0	1	1	9.84	12.120	
5445	2012/12/19 5:00	4	0	1	1	10.66	14.395	

5446	2012/12/19 6:00	4	0	1	1	9.84	12.880
5447	2012/12/19 7:00	4	0	1	1	10.66	13.635
5448	2012/12/19 8:00	4	0	1	1	9.84	12.880
5449	2012/12/19 9:00	4	0	1	1	11.48	14.395
5450	2012/12/19 10:00	4	0	1	1	13.12	16.665
5451	2012/12/19 11:00	4	0	1	1	16.40	20.455
5452	2012/12/19 12:00	4	0	1	1	16.40	20.455
5453	2012/12/19 13:00	4	0	1	1	17.22	21.210
5454	2012/12/19 14:00	4	0	1	1	17.22	21.210
5455	2012/12/19 15:00	4	0	1	1	17.22	21.210
5456	2012/12/19 16:00	4	0	1	1	17.22	21.210
5457	2012/12/19 17:00	4	0	1	1	16.40	20.455
5458	2012/12/19 18:00	4	0	1	1	15.58	19.695
5459	2012/12/19 19:00	4	0	1	1	15.58	19.695
5460	2012/12/19 20:00	4	0	1	1	14.76	17.425
5461	2012/12/19 21:00	4	0	1	1	13.94	15.910
5462	2012/12/19 22:00	4	0	1	1	13.94	17.425
5463	2012/12/19 23:00	4	0	1	1	13.12	16.665

	humidity	windspeed	casual	registered	count
0	81	0.0000	3	13	16
1	80	0.0000	8	32	40
2	80	0.0000	5	27	32
3	75	0.0000	3	10	13
4	75	0.0000	0	1	1
5	75	6.0032	0	1	1
6	80	0.0000	2	0	2
7	86	0.0000	1	2	3
8	75	0.0000	1	7	8
9	76	0.0000	8	6	14
10	76	16.9979	12	24	36
11	81	19.0012	26	30	56
12	77	19.0012	29	55	84
13	72	19.9995	47	47	94
14	72	19.0012	35	71	106
15	77	19.9995	40	70	110
16	82	19.9995	41	52	93
17	82	19.0012	15	52	67
18	88	16.9979	9	26	35
19	88	16.9979	6	31	37
20	87	16.9979	11	25	36
21	87	12.9980	3	31	34
22	94	15.0013	11	17	28
23	88	19.9995	15	24	39
24	88	19.9995	4	13	17
25	94	16.9979	1	16	17
26	100	19.0012	1	8	9
27	94	12.9980	2	4	6

28	94	12.9980	2	1	3
29	77	19.9995	0	2	2
...
5434	46	22.0028	13	512	525
5435	46	26.0027	19	334	353
5436	50	16.9979	4	264	268
5437	50	15.0013	9	159	168
5438	49	0.0000	5	127	132
5439	49	6.0032	1	80	81
5440	61	0.0000	6	35	41
5441	65	6.0032	1	14	15
5442	65	6.0032	1	2	3
5443	75	8.9981	0	5	5
5444	75	8.9981	1	6	7
5445	75	6.0032	2	29	31
5446	75	6.0032	3	109	112
5447	75	8.9981	3	360	363
5448	87	7.0015	13	665	678
5449	75	7.0015	8	309	317
5450	70	7.0015	17	147	164
5451	54	15.0013	31	169	200
5452	54	19.0012	33	203	236
5453	50	12.9980	30	183	213
5454	50	12.9980	33	185	218
5455	50	19.0012	28	209	237
5456	50	23.9994	37	297	334
5457	50	26.0027	26	536	562
5458	50	23.9994	23	546	569
5459	50	26.0027	7	329	336
5460	57	15.0013	10	231	241
5461	61	15.0013	4	164	168
5462	61	6.0032	12	117	129
5463	66	8.9981	4	84	88

[10886 rows x 12 columns]

Then we let pandas to tell us some information we have to know features name and type at the beginning

In [61]: df.dtypes

```
Out[61]: datetime    object
season             int64
holiday            int64
workingday         int64
weather            int64
temp              float64
atemp             float64
```

```

humidity      int64
windspeed     float64
casual         int64
registered     int64
count         int64
dtype: object

```

then we should know how large the dataset is

```
In [62]: df.shape
```

```
Out[62]: (10886, 12)
```

In conclusion we have 10886 rows each row has 12 different features. Also, there might be some noise data to deal with, so let's see if there are some missing values.

```
In [63]: df.count()
```

```

Out[63]: datetime      10886
season                10886
holiday              10886
workingday          10886
weather             10886
temp               10886
atemp             10886
humidity           10886
windspeed          10886
casual             10886
registered          10886
count              10886
dtype: int64

```

we can see that there is no missing values

```
In [65]: type(df.datetime)
```

```
Out[65]: pandas.core.series.Series
```

Let's process the time feature, since it has much more information and target value always changes with time.

```

In [66]: df['month'] = pd.DatetimeIndex(df.datetime).month
df['day'] = pd.DatetimeIndex(df.datetime).dayofweek
df['hour'] = pd.DatetimeIndex(df.datetime).hour

```

```
In [67]: df.head(10)
```

```

Out[67]:
   datetime      season  holiday  workingday  weather  temp  atemp  \
0  2011/1/1 0:00        1        0           0        1   9.84  14.395
1  2011/1/1 1:00        1        0           0        1   9.02  13.635

```

2	2011/1/1 2:00	1	0	0	1	9.02	13.635
3	2011/1/1 3:00	1	0	0	1	9.84	14.395
4	2011/1/1 4:00	1	0	0	1	9.84	14.395
5	2011/1/1 5:00	1	0	0	2	9.84	12.880
6	2011/1/1 6:00	1	0	0	1	9.02	13.635
7	2011/1/1 7:00	1	0	0	1	8.20	12.880
8	2011/1/1 8:00	1	0	0	1	9.84	14.395
9	2011/1/1 9:00	1	0	0	1	13.12	17.425

	humidity	windspeed	casual	registered	count	month	day	hour
0	81	0.0000	3	13	16	1	5	0
1	80	0.0000	8	32	40	1	5	1
2	80	0.0000	5	27	32	1	5	2
3	75	0.0000	3	10	13	1	5	3
4	75	0.0000	0	1	1	1	5	4
5	75	6.0032	0	1	1	1	5	5
6	80	0.0000	2	0	2	1	5	6
7	86	0.0000	1	2	3	1	5	7
8	75	0.0000	1	7	8	1	5	8
9	76	0.0000	8	6	14	1	5	9

After preprocessing time series feature, we can drop original time features And in this base-line version, we don't use registered feature as well

```
In [68]: df_origin = df
         df = df.drop(['datetime', 'casual', 'registered'], axis = 1)
```

```
In [69]: df.head(5)
```

```
Out [69]:
```

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	\
0	1	0	0	1	9.84	14.395	81	0.0	
1	1	0	0	1	9.02	13.635	80	0.0	
2	1	0	0	1	9.02	13.635	80	0.0	
3	1	0	0	1	9.84	14.395	75	0.0	
4	1	0	0	1	9.84	14.395	75	0.0	

	count	month	day	hour
0	16	1	5	0
1	40	1	5	1
2	32	1	5	2
3	13	1	5	3
4	1	1	5	4

Well, that seems more clear

```
In [70]: df.shape
```

```
Out [70]: (10886, 12)
```

separate dataset into two: 1. df_targetgoalcount feature 2. df_datadata

```
In [73]: df_target = df['count'].values
         df_data = df.drop(['count'],axis = 1).values
         print('df_data shape is ', df_data.shape)
         print('df_target shape is ', df_target.shape)
```

```
df_data shape is (10886, 11)
df_target shape is (10886,)
```

3 Machine Learning Algorithms

the process below shows that we might spend lots of time on parameter modification different parameters would lead to different results

```
In [74]: from sklearn import linear_model
         from sklearn import cross_validation
         from sklearn import svm
         from sklearn.ensemble import RandomForestRegressor
         from sklearn.learning_curve import learning_curve
         from sklearn.grid_search import GridSearchCV
         from sklearn.metrics import explained_variance_score
```

```
/Users/songzhewei/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning:
    "This module will be removed in 0.20.", DeprecationWarning)
/Users/songzhewei/anaconda3/lib/python3.6/site-packages/sklearn/learning_curve.py:22: DeprecationWarning:
    DeprecationWarning)
/Users/songzhewei/anaconda3/lib/python3.6/site-packages/sklearn/grid_search.py:42: DeprecationWarning:
    DeprecationWarning)
```

data size is small, we are going to try different algorithms We would use cross validation validation data is 20% to see model's performance we would try Support Vector Regression, Ridge Regression and Random Forest Regressor

```
In [84]: cv = cross_validation.ShuffleSplit(len(df_data), n_iter=3, test_size=0.2,
         random_state=0)

         print("ridge")
         for train, test in cv:
             svc = linear_model.Ridge().fit(df_data[train], df_target[train])
             print("train score: {0:.3f}, test score: {1:.3f}\n".format(
                 svc.score(df_data[train], df_target[train]), svc.score(df_data[test], df_target[test]))

         print("SVR(kernel='rbf',C=10,gamma=.001)")
         for train, test in cv:

             svc = svm.SVR(kernel='rbf', C = 10, gamma = .001).fit(df_data[train], df_target[train])
             print("train score: {0:.3f}, test score: {1:.3f}\n".format(
```

```

        svc.score(df_data[train], df_target[train]), svc.score(df_data[test], df_target[test])

print ("Random Forest(n_estimators = 100)")
for train, test in cv:
    svc = RandomForestRegressor(n_estimators = 100).fit(df_data[train], df_target[train])
    print("train score: {0:.3f}, test score: {1:.3f}\n".format(
        svc.score(df_data[train], df_target[train]), svc.score(df_data[test], df_target[test])

ridge
train score: 0.339, test score: 0.332

train score: 0.330, test score: 0.370

train score: 0.342, test score: 0.320

SVR(kernel='rbf',C=10,gamma=.001)
train score: 0.417, test score: 0.408

train score: 0.406, test score: 0.452

train score: 0.419, test score: 0.390

Random Forest(n_estimators = 100)
train score: 0.982, test score: 0.866

train score: 0.981, test score: 0.881

train score: 0.982, test score: 0.868

```

Random forest has best performance Next step we are going to do parameter modification
There is tools call grid search which can help us find optimal parameter

```

In [88]: X = df_data
        y = df_target

X_train, X_test, y_train, y_test = cross_validation.train_test_split(
    X, y, test_size=0.2, random_state=0)

tuned_parameters = [{'n_estimators': [10, 100, 500]}]

scores = ['r2']

for score in scores:

    print(score)

```



```

clf = GridSearchCV(RandomForestRegressor(), tuned_parameters, cv=5, scoring=score)
clf.fit(X_train, y_train)

print("we found optimal parameter")
print( "")
#best_estimator_ returns the best estimator chosen by the search
print(clf.best_estimator_)
print( "")
print("score is:")
print( "")

for params, mean_score, scores in clf.grid_scores_:
    print("%0.3f (+/-%0.03f) for %r"
          % (mean_score, scores.std() / 2, params))
print( "")

```

r2

we found optimal parameter

```

RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
    max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=500, n_jobs=1,
    oob_score=False, random_state=None, verbose=0, warm_start=False)

```

score is:

```

0.847 (+/-0.008) for {'n_estimators': 10}
0.862 (+/-0.006) for {'n_estimators': 100}
0.863 (+/-0.006) for {'n_estimators': 500}

```

We can see Grid Search is helpful we use these parameter on our model we also need to check whether our model is overfitting plot learning curve

```

In [89]: def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
    n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):

    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

```

```

train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt


```

```

title = "Learning Curves (Random Forest, n_estimators = 100)"
cv = cross_validation.ShuffleSplit(df_data.shape[0], n_iter=10, test_size=0.2, random_state=0)
estimator = RandomForestRegressor(n_estimators = 100)
plot_learning_curve(estimator, title, X, y, (0.0, 1.01), cv=cv, n_jobs=4)

plt.show()

```



There is a big gap between training curve and test curve, overfitting occurred

```
In [25]: # mitigate overfitting
print "Random Forest(n_estimators=200, max_features=0.6, max_depth=15)"
for train, test in cv:
    svc = RandomForestRegressor(n_estimators = 200, max_features=0.6, max_depth=15).fit(
    print("train score: {0:.3f}, test score: {1:.3f}\n".format(
        svc.score(df_data[train], df_target[train]), svc.score(df_data[test], df_target[test])

/Random Forest(n_estimators=200, max_features=0.3)
train score: 0.965, test score: 0.867

train score: 0.966, test score: 0.885

train score: 0.966, test score: 0.875

train score: 0.965, test score: 0.876

train score: 0.967, test score: 0.870

train score: 0.965, test score: 0.872

train score: 0.967, test score: 0.862

train score: 0.966, test score: 0.875

train score: 0.966, test score: 0.871

train score: 0.966, test score: 0.868
```

we can use registered feature to do prediction separate dataset into two

```
In [26]: df_registered = df_origin.drop(['datetime', 'casual', 'count'], axis = 1)
df_casual = df_origin.drop(['datetime', 'count', 'registered'], axis = 1)
```

```
In [27]: df_train_registered.head()
```

```
Out[27]:
```

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	\
0	1	0	0	1	9.84	14.395	81	0.0	
1	1	0	0	1	9.02	13.635	80	0.0	
2	1	0	0	1	9.02	13.635	80	0.0	
3	1	0	0	1	9.84	14.395	75	0.0	
4	1	0	0	1	9.84	14.395	75	0.0	

registered month day hour

0	13	1	5	0
1	32	1	5	1
2	27	1	5	2
3	10	1	5	3
4	1	1	5	4

In [29]: df_train_casual.head()

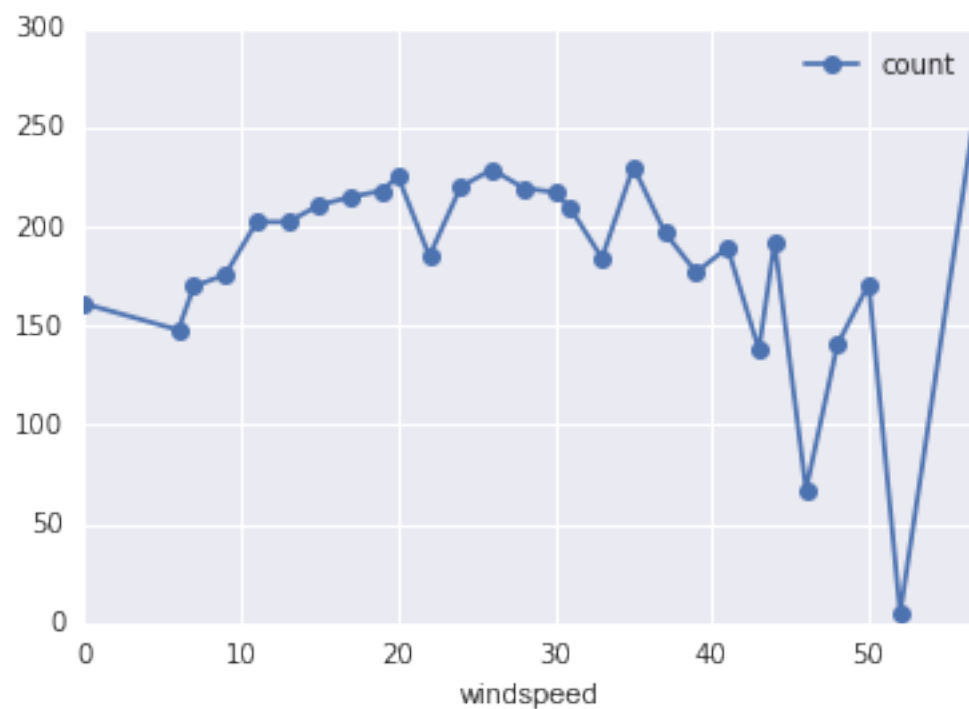
Out[29]:

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	\
0	1	0	0	1	9.84	14.395	81	0.0	
1	1	0	0	1	9.02	13.635	80	0.0	
2	1	0	0	1	9.02	13.635	80	0.0	
3	1	0	0	1	9.84	14.395	75	0.0	
4	1	0	0	1	9.84	14.395	75	0.0	

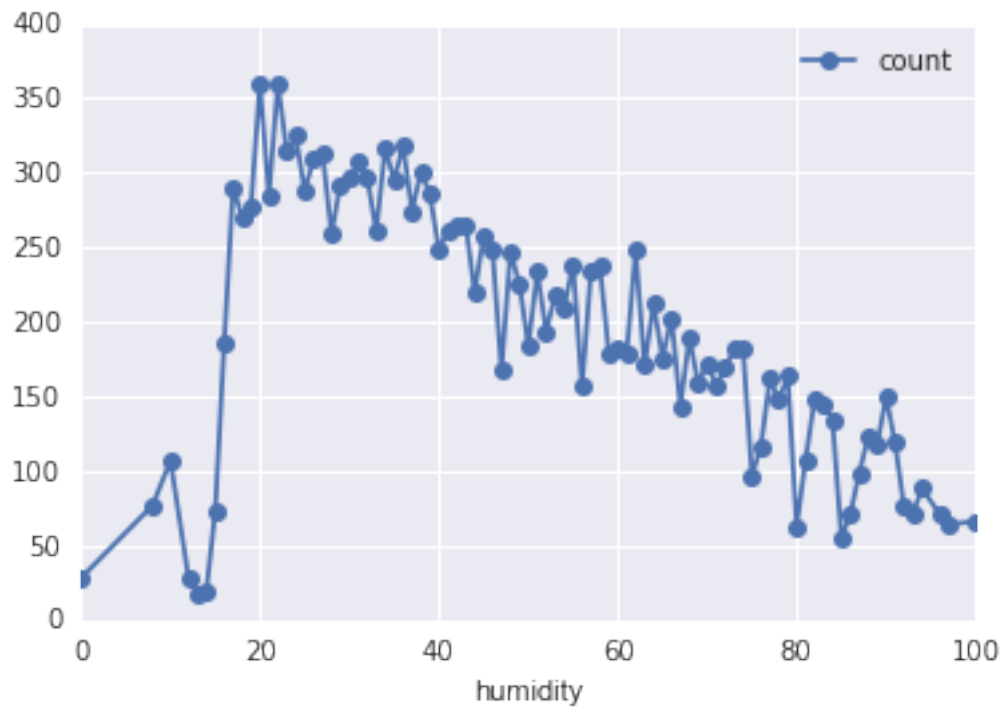
	casual	month	day	hour
0	3	1	5	0
1	8	1	5	1
2	5	1	5	2
3	3	1	5	3
4	0	1	5	4

Data analysis and visulization

In [40]: # *windspeed*
df_origin.groupby('windspeed').mean().plot(y='count', marker='o')
plt.show()




```
In [41]: # humidity
df_origin.groupby('humidity').mean().plot(y='count', marker='o')
plt.show()
```



```
In [42]: # temperature
df_origin.groupby('temp').mean().plot(y='count', marker='o')
plt.show()
```

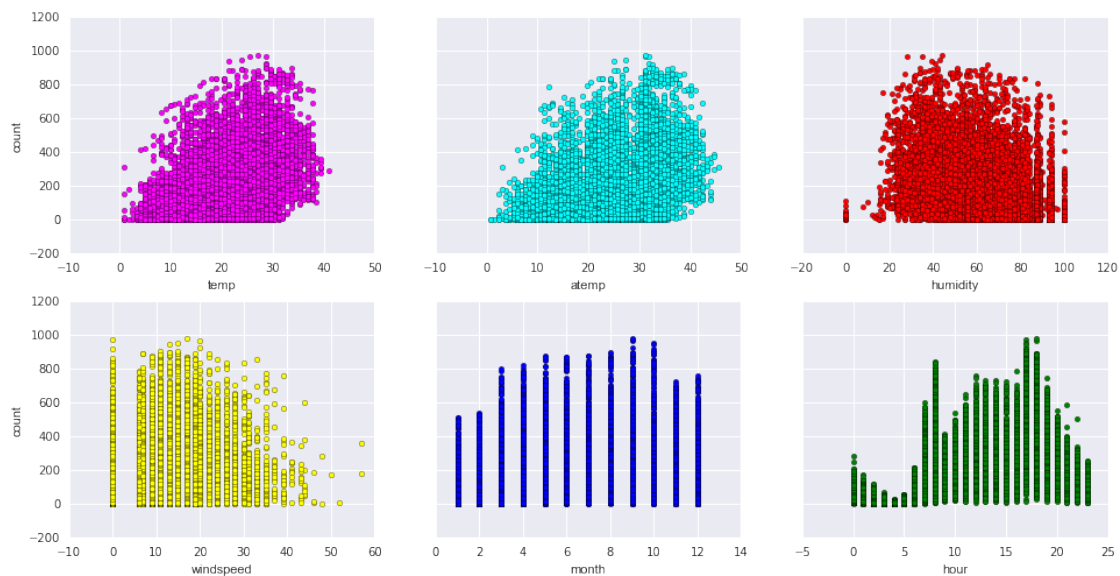


```
In [46]: #temp humidity changing
df_train_origin.plot(x='temp', y='humidity', kind='scatter')
plt.show()
```




```
In [35]: # scatter different dimentions distribution
fig, axs = plt.subplots(2, 3, sharey=True)
df_origin.plot(kind='scatter', x='temp', y='count', ax=axs[0, 0], figsize=(16, 8), color='magenta')
df_origin.plot(kind='scatter', x='atemp', y='count', ax=axs[0, 1], color='cyan')
df_origin.plot(kind='scatter', x='humidity', y='count', ax=axs[0, 2], color='red')
df_origin.plot(kind='scatter', x='windspeed', y='count', ax=axs[1, 0], color='yellow')
df_origin.plot(kind='scatter', x='month', y='count', ax=axs[1, 1], color='blue')
df_origin.plot(kind='scatter', x='hour', y='count', ax=axs[1, 2], color='green')
```

Out [35]: <matplotlib.axes._subplots.AxesSubplot at 0x11ad48090>



```
In [37]: sns.pairplot(df_origin[["temp", "month", "humidity", "count"]], hue="count")
```

Out [37]: <seaborn.axisgrid.PairGrid at 0x11beec90>




```
In [48]: # correlation analysis
corr = df_origin[['temp', 'weather', 'windspeed', 'day', 'month', 'hour', 'count']].corr()
corr
```

```
Out[48]:
```

	temp	weather	windspeed	day	month	hour	\
temp	1.000000	-0.055035	-0.017852	-0.038466	0.257589	0.145430	
weather	-0.055035	1.000000	0.007261	-0.047692	0.012144	-0.022740	
windspeed	-0.017852	0.007261	1.000000	-0.024804	-0.150192	0.146631	
day	-0.038466	-0.047692	-0.024804	1.000000	-0.002266	-0.002925	
month	0.257589	0.012144	-0.150192	-0.002266	1.000000	-0.006818	
hour	0.145430	-0.022740	0.146631	-0.002925	-0.006818	1.000000	
count	0.394454	-0.128655	0.101369	-0.002283	0.166862	0.400601	

	count
temp	0.394454
weather	-0.128655
windspeed	0.101369
day	-0.002283
month	0.166862
hour	0.400601
count	1.000000

```
In [52]: plt.figure()
plt.matshow(corr)
plt.colorbar()
plt.show()
```

<matplotlib.figure.Figure at 0x14716d410>

